

ML@NOVA DFJ

Predicting Survival Time in Multiple Myeloma Patients

Team identification

Name 1: José Costa

Number 1: 72899

Name 2: Dinis Santos

Number 2: 72815

Name 3: Francisco Jorge

Number 3: 70293

Final score on the private leaderboard:

Leaderboard private ranking:

Task [1.1] - Data preparation and validation pipeline

What was done in task [1.1]

1. Missing Values Analysis

- Visualized missing values using multiple methods (bar plot, heatmap, matrix, dendrogram)
- Created comprehensive overview of data completeness
- Identified patterns in missing data

2. Data Cleaning

- Dropped rows with missing SurvivalTime values
- Removed columns containing missing data (baseline approach)
- Excluded censored cases (where Censored == 1)
- Retained only complete, uncensored observations

3. Feature Exploration

- Visualized feature relationships using pairplot
- Analyzed correlations between Age, Gender, Stage, TreatmentType, and SurvivalTime
- Examined distribution patterns across features

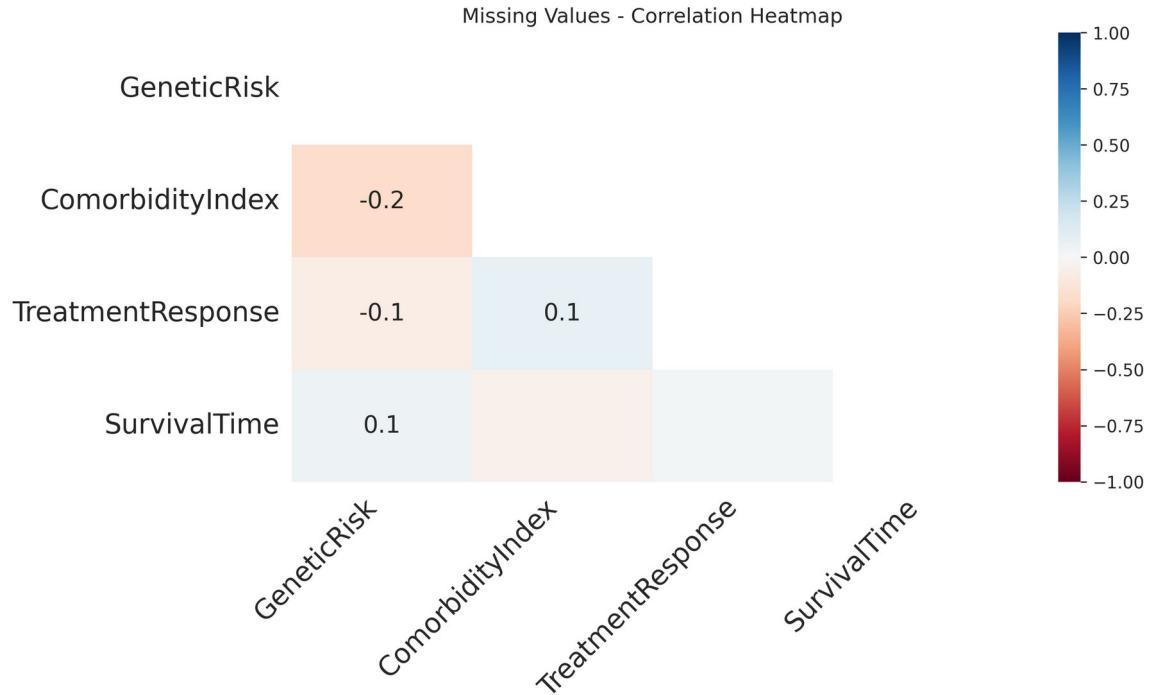
What was done in task [1.1]

| | |
|------------------------------------|---|
| 4. Data Preparation | Defined feature matrix (X) by dropping target and identifier columns |
| | Isolated target variable (y) as SurvivalTime |
| | Preserved censoring indicator for potential future use |
| 5. Validation Strategy Development | Implemented train/validation/test split (64%/16%/20%) |
| | Tested simple split approach with Linear Regression |
| | Implemented 5-fold cross-validation for more robust evaluation |
| | Compared both validation strategies (simple split vs. cross-validation) |
| 6. Performance Evaluation | Calculated MSE (Mean Squared Error) and cMSE (Censored MSE) |
| | Evaluated model performance on validation and test sets |
| | Compared cross-validation results to simple split results |

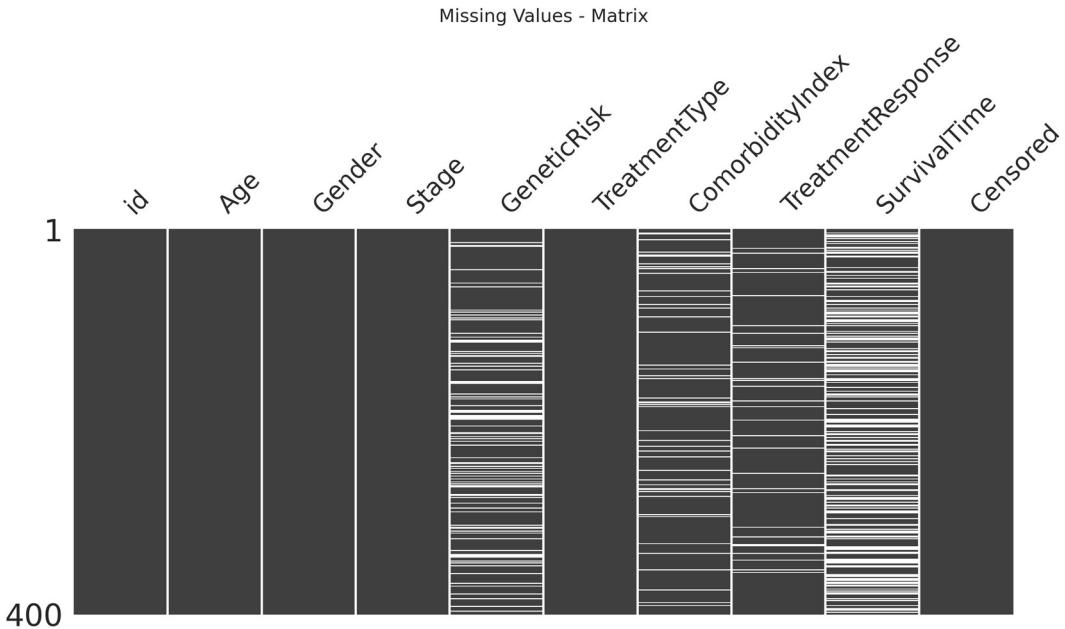
Results and Analysis from task [1.1]



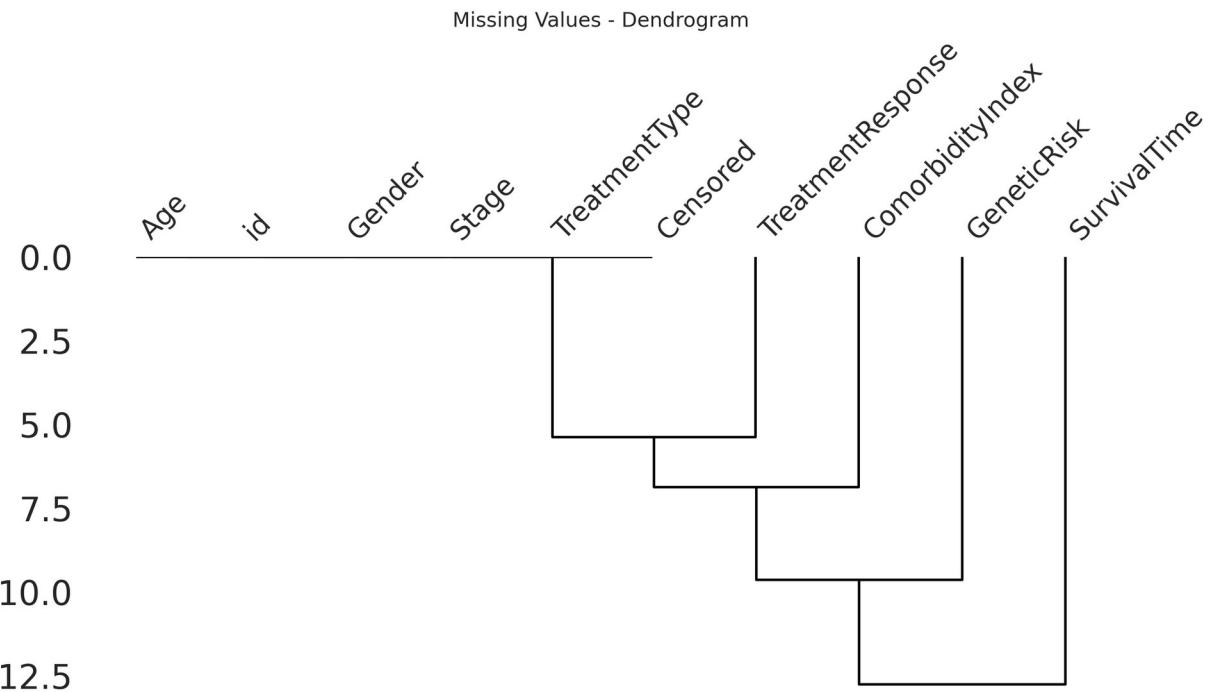
Results and Analysis from task [1.1]



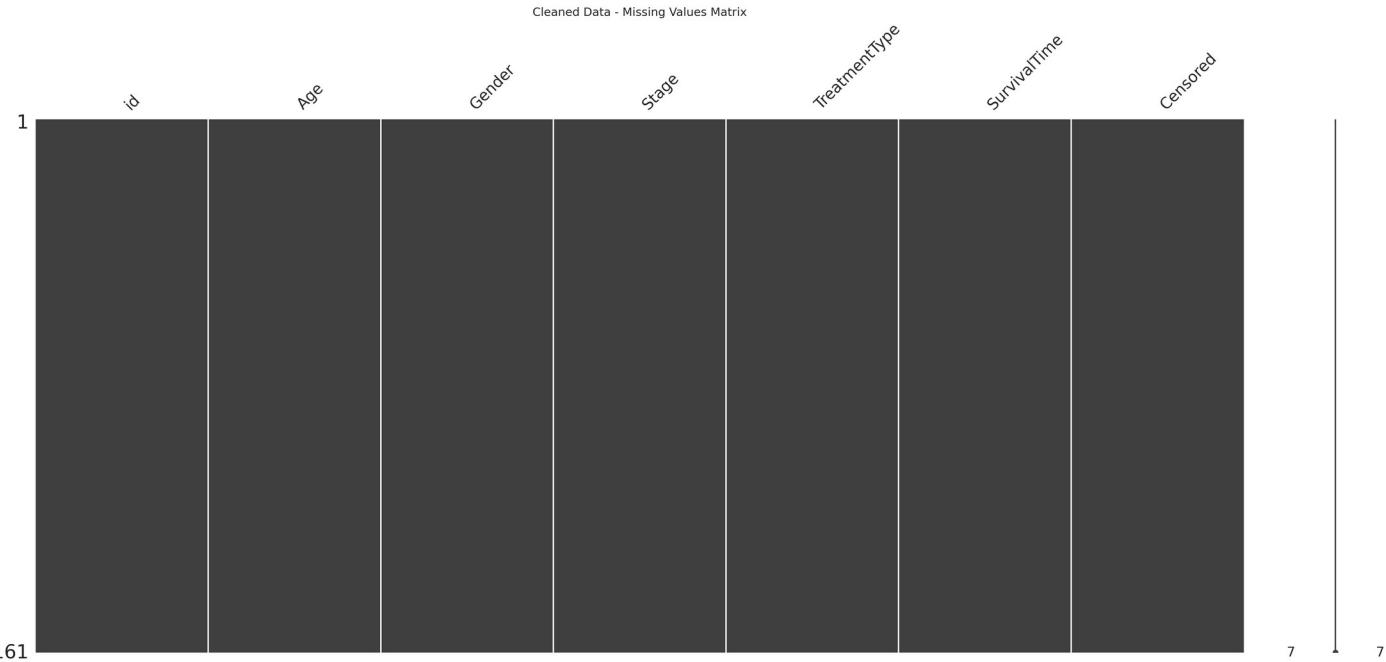
Results and Analysis from task [1.1]



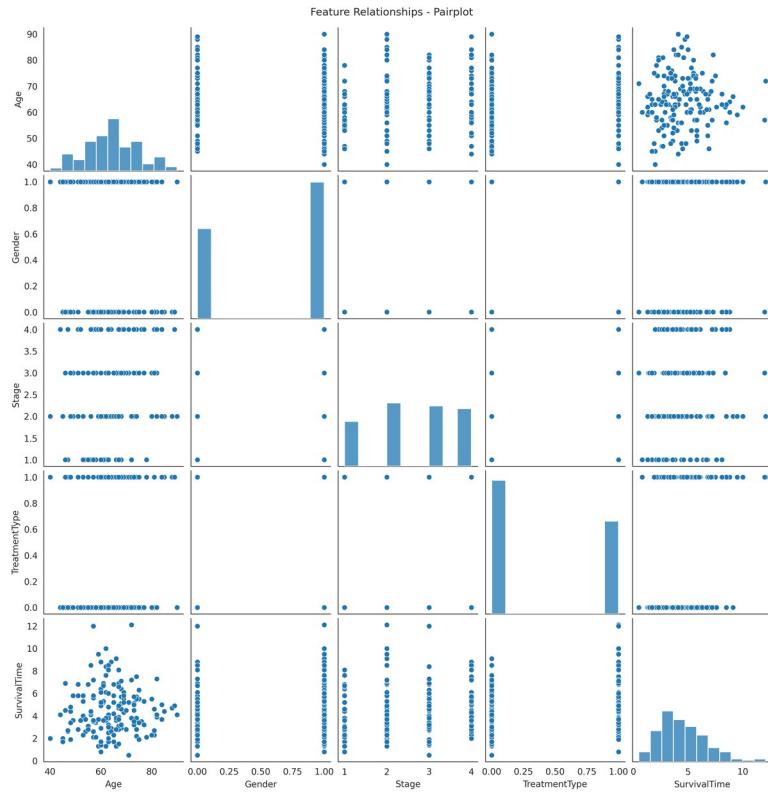
Results and Analysis from task [1.1]



Results and Analysis from task [1.1]



Results and Analysis from task [1.1]



Task [1.2] - Learn the baseline model

What was done in task [1.2]

| | |
|------------------------------|--|
| 1. Pipeline Construction | - Built baseline pipeline combining StandardScaler and Linear Regression - Ensured feature scaling for improved model performance - Created modular, reusable pipeline structure |
| 2. Cross-Validation Training | - Performed 5-fold cross-validation for robust model evaluation - Calculated CV MSE scores across all folds - Computed mean and standard deviation of cross-validation performance |
| 3. Final Model Training | - Fitted baseline pipeline on entire training dataset - Generated predictions on training data - Maximized use of available data for final model |

What was done in task [1.2]

4.

Performance Metrics

- Calculated Training MSE (Mean Squared Error)

- Calculated Training cMSE (Censored Mean Squared Error)

- Established baseline performance benchmarks

5. Test

Predictions & Submission

- Loaded test dataset and prepared features

- Generated predictions for test samples

- Created submission file for competition/evaluation

6. Model

Visualization

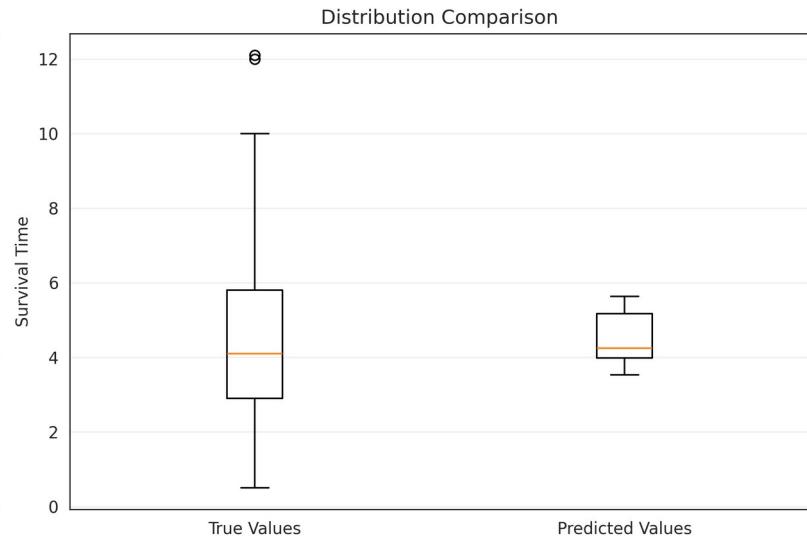
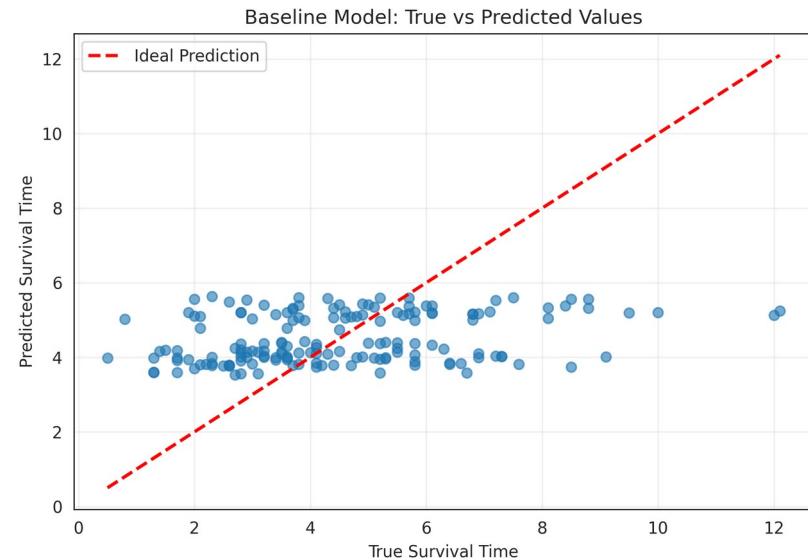
- Created scatter plot comparing true vs predicted survival times

- Generated boxplot for distribution comparison

- Visualized model fit quality and prediction patterns

- Saved individual plots for documentation

Results and Analysis from task [1.2]



Task [2.1] - Development

What was done in task [2.1]

| | |
|---|--|
| 1. Polynomial Regression Function Development | - Created `train_polynomial_regression()` function with hyperparameter search - Implemented cross-validation for degree selection (testing degrees 1 to max_degree) - Added early stopping mechanism (stops after 2 consecutive iterations without improvement) - Returned best degree, trained model, and complete CV results dictionary |
| 2. k-Nearest Neighbors Function Development | - Created `train_knn()` function with hyperparameter search - Implemented cross-validation for k selection (testing k from 1 to max_k) - Added early stopping mechanism for efficiency - Returned best k value, trained model, and complete CV results dictionary |
| 3. Hyperparameter Selection | - Used 5-fold cross-validation for both models - Searched polynomial degrees from 1 to 10 - Searched k values from 1 to 20 - Tracked MSE scores with standard deviations for each hyperparameter |

What was done in task [2.1]

4. Model Training

- Trained Polynomial Regression with optimal degree on full dataset
 - Trained k-NN Regression with optimal k on full dataset
 - Generated predictions on training data for both models
-

5. Performance Evaluation

- Calculated training MSE for both models
 - Calculated training cMSE for both models
 - Compared performance against baseline expectations
 - Documented hyperparameter selection results
-

Results and Analysis from task [2.1]

Task [2.2] - Evaluation

What was done in task [2.2]

1. Comprehensive Model Comparison

- Created comparison table with baseline, polynomial regression, and k-NN models
- Included hyperparameter configurations for each model
- Displayed min, max, mean, and standard deviation of errors
- Identified best performing model based on mean cross-validation error

2. Hyperparameter Tuning Visualization

- Plotted polynomial degree vs MSE with confidence intervals
- Plotted k-value vs MSE with confidence intervals
- Marked optimal hyperparameters with vertical lines
- Showed performance trends across hyperparameter ranges

3. Model Predictions Comparison

- Created scatter plots of true vs predicted values for all three models
- Displayed MSE on each plot for direct comparison
- Included ideal prediction line ($y=x$) as reference
- Generated combined and individual visualization plots

What was done in task [2.2]

4. Statistical Analysis

- Computed cross-validation statistics for each model
- Analyzed variance in predictions across folds
- Compared model stability through standard deviation metrics
- Evaluated improvement over baseline model

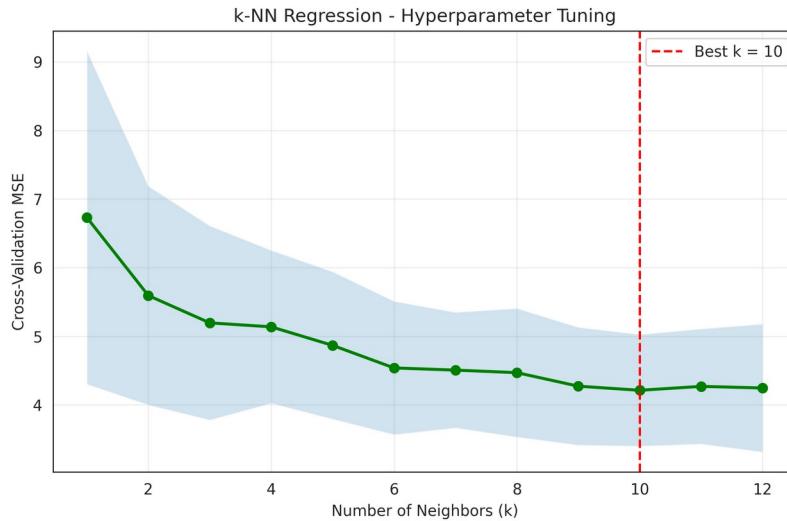
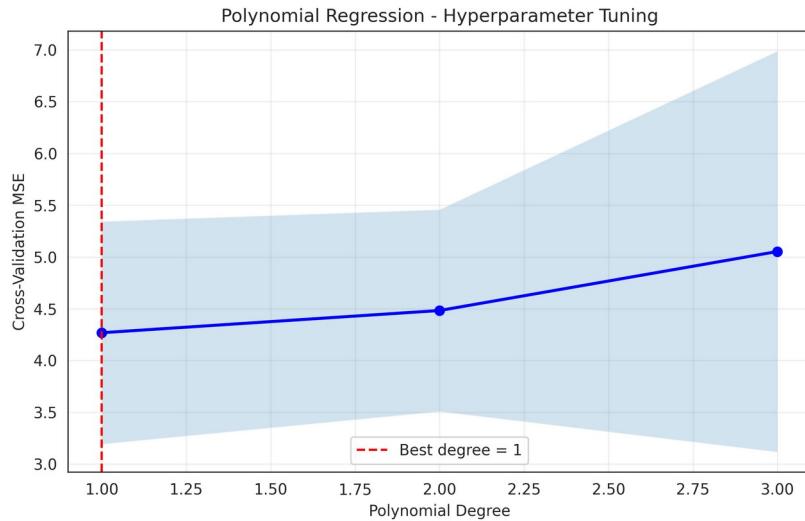
5. Test Set Predictions

- Selected best performing model based on CV results
- Generated predictions for test dataset
- Created submission file for evaluation
- Documented model selection rationale

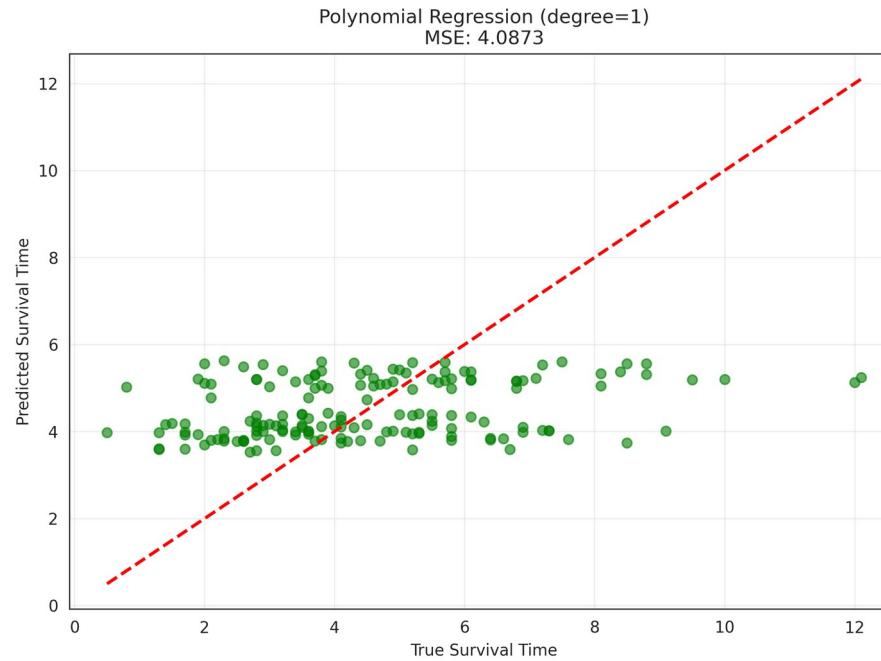
6. Results Documentation

- Saved all comparison plots with task-specific naming
- Generated separate plots for polynomial and k-NN tuning
- Created individual prediction visualizations for each model
- Documented complete evaluation workflow

Results and Analysis from task [2.2]

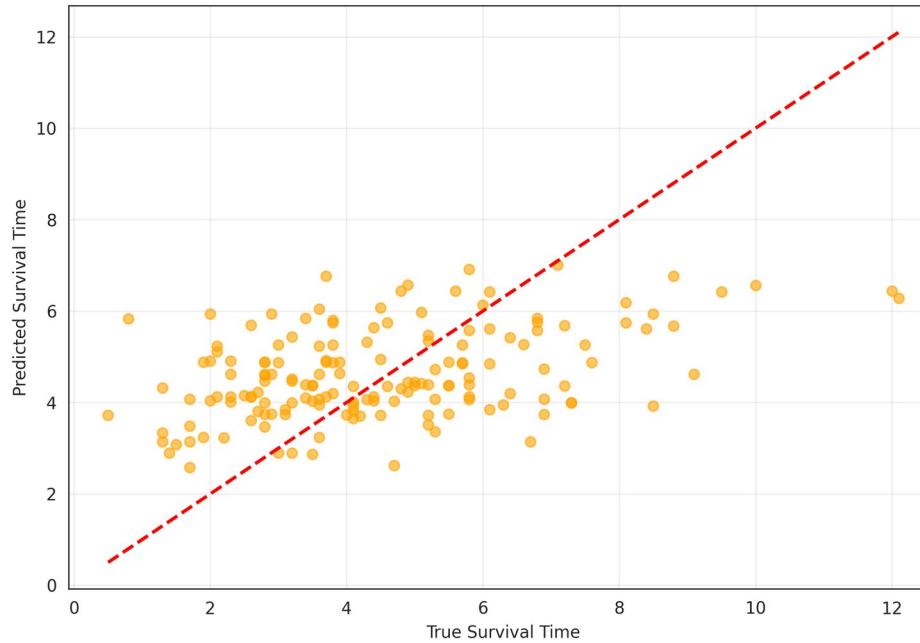


Results and Analysis from task [2.2]

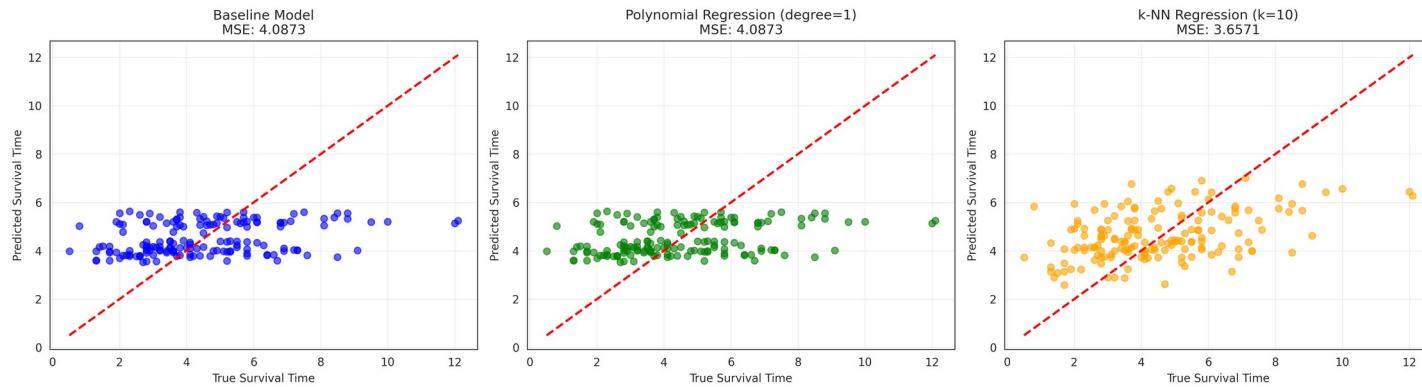


Results and Analysis from task [2.2]

k-NN Regression (k=10)
MSE: 3.6571



Results and Analysis from task [2.2]



Code Demo

Overall assessment

What went wrong

What went great