

ML@NOVA DFJ

Predicting Survival Time in Multiple Myeloma Patients

Team identification

Name 1: José Costa

Number 1: 72899

Name 2: Dinis Santos

Number 2: 72815

Name 3: Francisco Jorge

Number 3: 70293

Final score on the private leaderboard: 2.58777

Leaderboard private ranking: 30

Task [1.1] - Data preparation and validation pipeline

What was done in task [1.1]

1. Missing Values Analysis

- Visualized missing values using multiple methods (bar plot, heatmap, matrix, dendrogram)
- Created comprehensive overview of data completeness
- Identified patterns in missing data

2. Data Cleaning

- Dropped rows with missing SurvivalTime values
- Removed columns containing missing data (baseline approach)
- Excluded censored cases (where Censored == 1)
- Retained only complete, uncensored observations

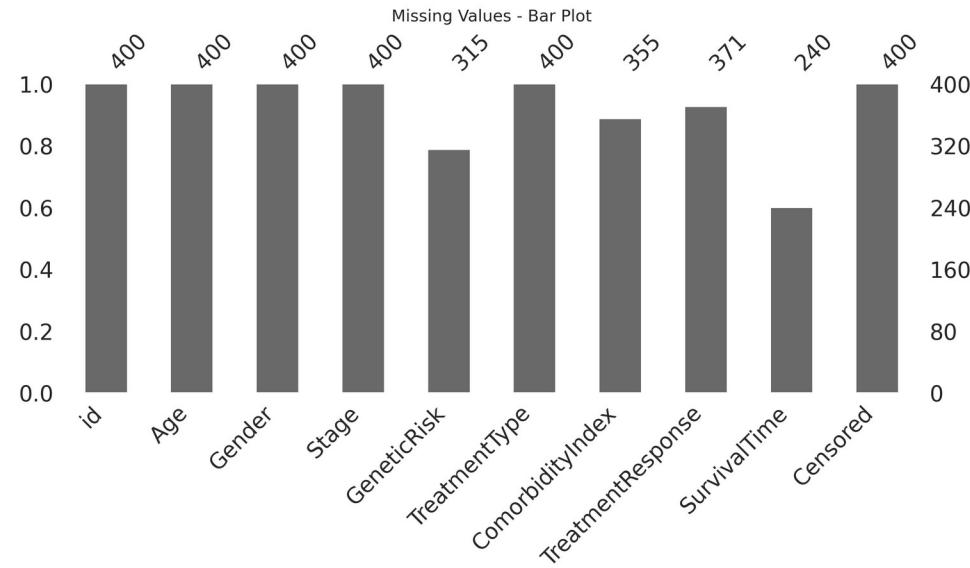
3. Feature Exploration

- Visualized feature relationships using pairplot
- Analyzed correlations between Age, Gender, Stage, TreatmentType, and SurvivalTime
- Examined distribution patterns across features

What was done in task [1.1]

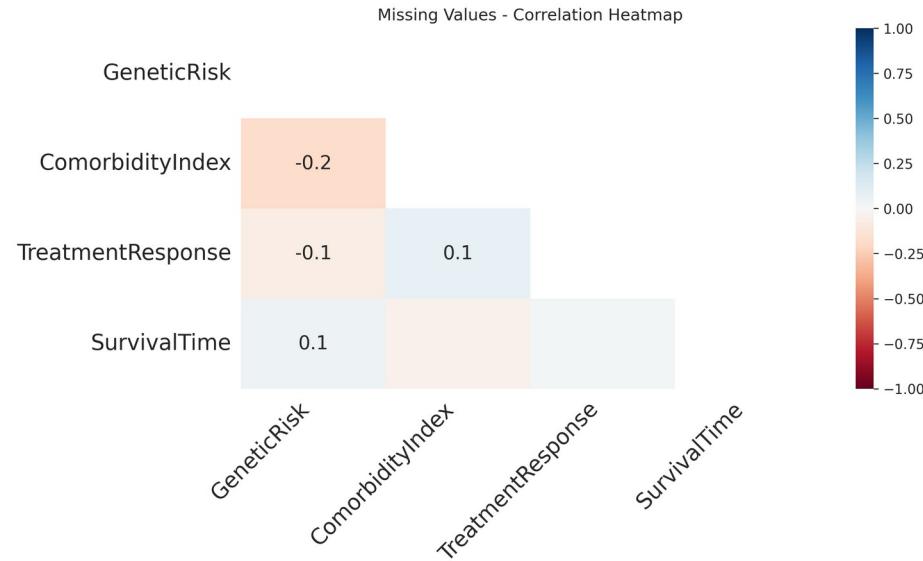
| | |
|------------------------------------|---|
| 4. Data Preparation | Defined feature matrix (X) by dropping target and identifier columns |
| | Isolated target variable (y) as SurvivalTime |
| | Preserved censoring indicator for potential future use |
| 5. Validation Strategy Development | Implemented train/validation/test split (64%/16%/20%) |
| | Tested simple split approach with Linear Regression |
| | Implemented 5-fold cross-validation for more robust evaluation |
| | Compared both validation strategies (simple split vs. cross-validation) |
| 6. Performance Evaluation | Calculated MSE (Mean Squared Error) and cMSE (Censored MSE) |
| | Evaluated model performance on validation and test sets |
| | Compared cross-validation results to simple split results |

Results and Analysis from task [1.1]



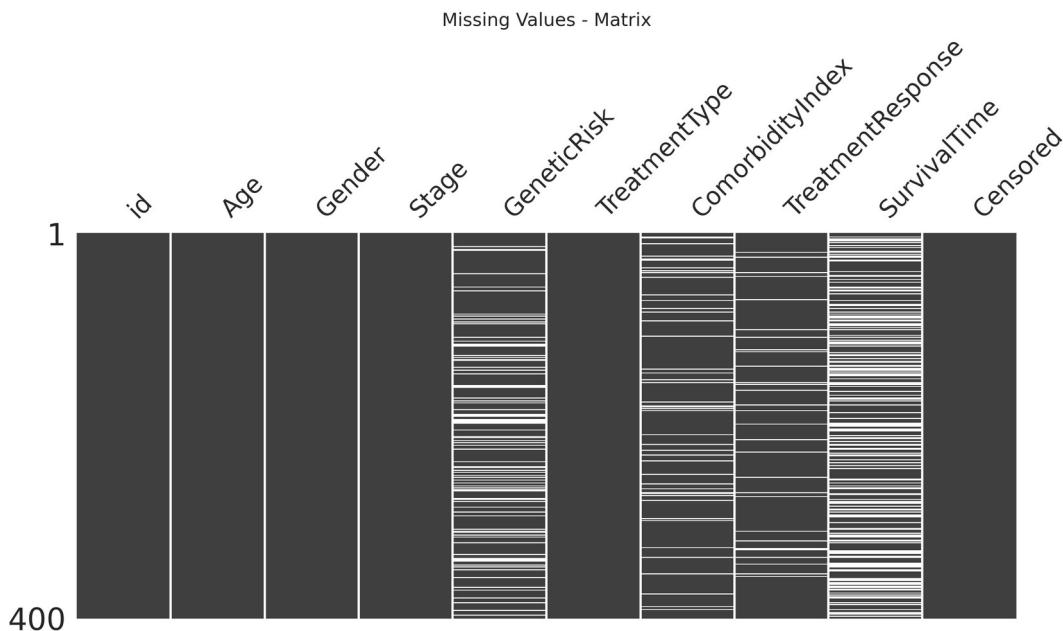
Bar chart showing missing-value counts per variable, with labels indicating how many entries are missing in each column.

Results and Analysis from task [1.1]



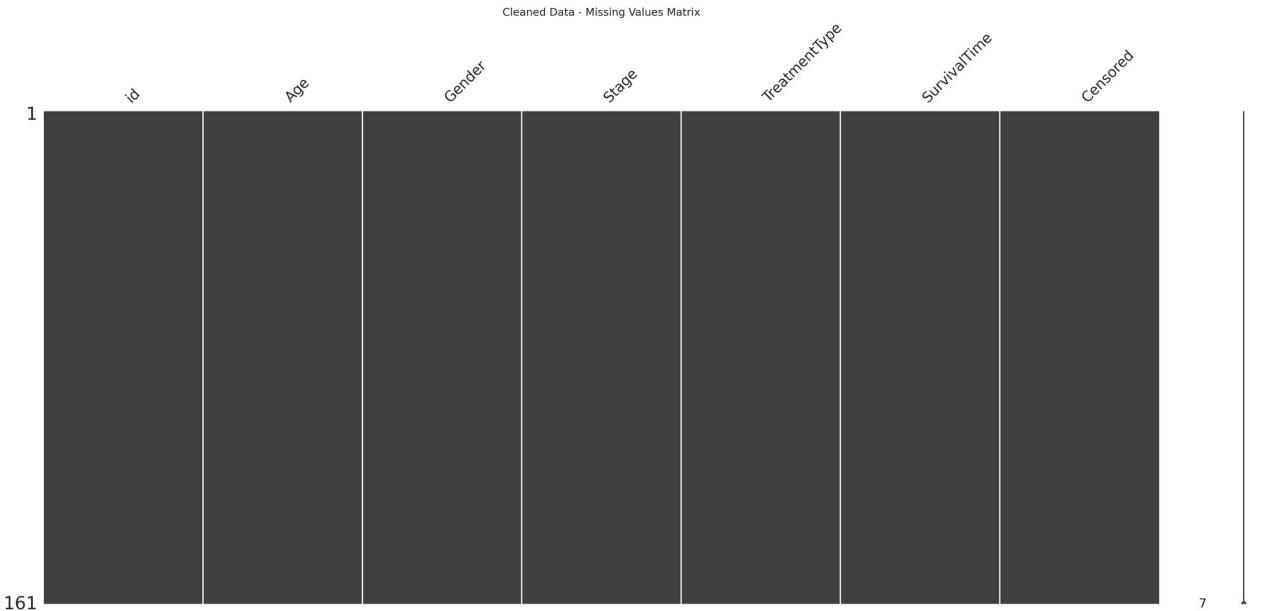
Heatmap showing correlations between missing-value patterns across selected variables, with mostly weak relationships.

Results and Analysis from task [1.1]



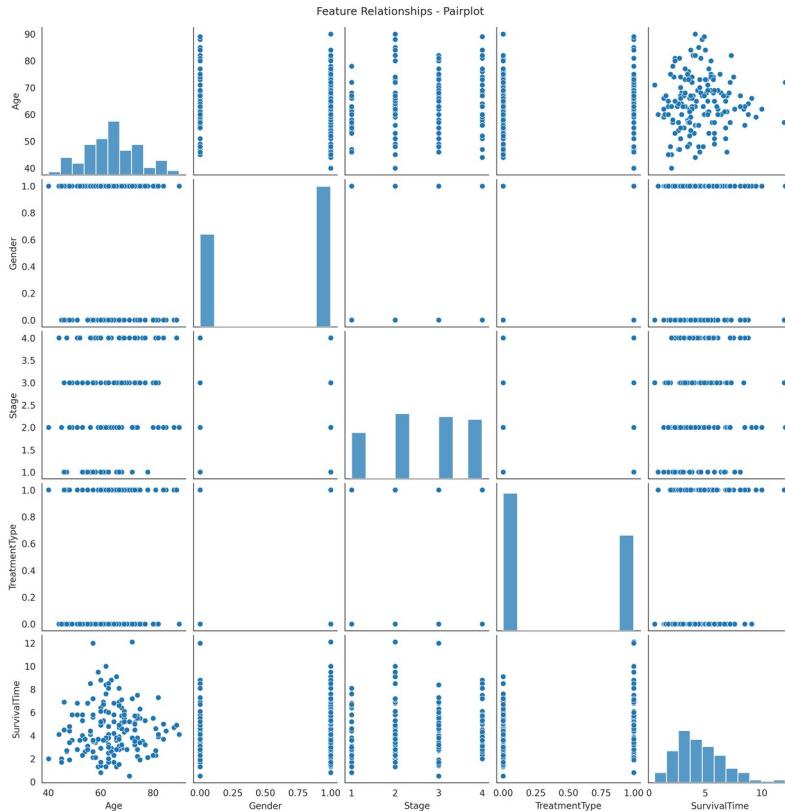
Matrix plot highlighting which entries are missing across variables, showing the distribution and density of gaps in the dataset.

Results and Analysis from task [1.1]



Matrix plot showing the cleaned dataset, confirming that all variables have no remaining missing values.

Results and Analysis from task [1.1]



- Age**
 - Roughly uniform distribution between 40 and 85.
 - No clear relationship with Stage, Gender, or TreatmentType.
 - Slight indication that lower ages may correspond to higher SurvivalTime, but with large variability.
- Gender**
 - Binary variable with no meaningful pattern across other features.
 - SurvivalTime similarly distributed between genders.
- Stage**
 - Balanced distribution across stages 1–4.
 - Clear trend: higher stages correspond to lower SurvivalTime.
 - No noticeable relationship with Age or Gender.
- TreatmentType**
 - Most individuals fall into one or two treatment categories.
 - No strong association with Age or Gender.
- SurvivalTime**
 - Right-skewed distribution, with most values between 1 and 5 years.
 - Strongest visible relationship is with Stage (more advanced stages → shorter survival).
 - No clear patterns with Gender or TreatmentType.

Results and Analysis from task [1.1]

| | MSE | cMSE |
|------------------|--------|--------|
| Simple Split | 5.0473 | 5.0473 |
| Cross-Validation | 4.0873 | 4.0873 |
| Pipeline | 4.4112 | 4.4112 |

The results show that cross-validation improves performance compared to a simple train–test split, reducing the MSE by almost one unit.

The pipeline approach performs slightly worse than cross-validation but still better than the simple split.

Task [1.2] - Learn the baseline model

What was done in task [1.2]

| | |
|------------------------------|--|
| 1. Pipeline Construction | - Built baseline pipeline combining StandardScaler and Linear Regression - Ensured feature scaling for improved model performance - Created modular, reusable pipeline structure |
| 2. Cross-Validation Training | - Performed 5-fold cross-validation for robust model evaluation - Calculated CV MSE scores across all folds - Computed mean and standard deviation of cross-validation performance |
| 3. Final Model Training | - Fitted baseline pipeline on entire training dataset - Generated predictions on training data - Maximized use of available data for final model |

What was done in task [1.2]

4.

Performance Metrics

- Calculated Training MSE (Mean Squared Error)

- Calculated Training cMSE (Censored Mean Squared Error)

- Established baseline performance benchmarks

5. Test

Predictions & Submission

- Loaded test dataset and prepared features

- Generated predictions for test samples

- Created submission file for competition/evaluation

6. Model

Visualization

- Created scatter plot comparing true vs predicted survival times

- Generated boxplot for distribution comparison

- Visualized model fit quality and prediction patterns

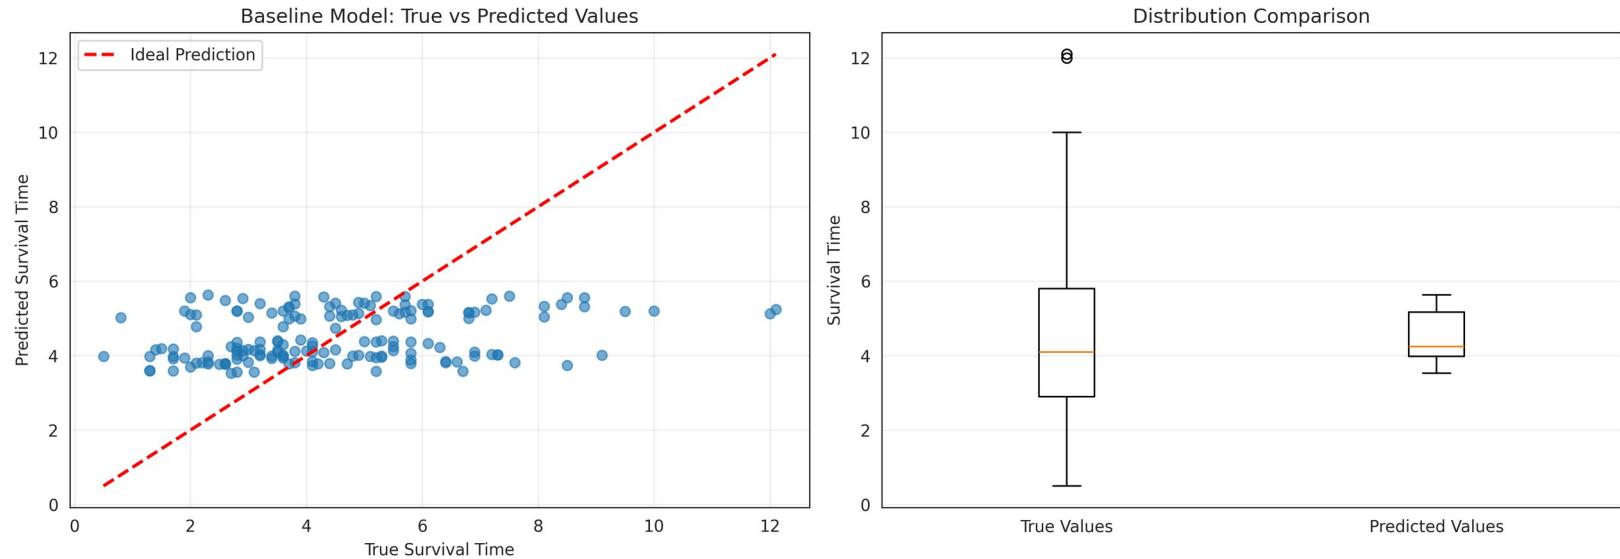
- Saved individual plots for documentation

Results and Analysis from task [1.2]

- Baseline Model Training with Cross-Validation

| MSE | cMSE | Best Kaggle Score |
|--------|--------|-------------------|
| 4.0873 | 4.0873 | 3.87347 |

Results and Analysis from task [1.2]



The scatter plot shows that the model's predictions are squeezed into a narrow band, mostly around 4–6 years and rarely follow the true values across their full range.

The boxplots make this even clearer: the true survival times vary widely, while the predicted ones stay tightly grouped.

Overall, the model isn't capturing the real variability in the data and is underfitting.

Task [2.1] - Development

What was done in task [2.1]

1. Polynomial Regression Function Development

- Created `train_polynomial_regression()` function with hyperparameter search
- Implemented cross-validation for degree selection (testing degrees 1 to max_degree)
- Added early stopping mechanism (stops after 2 consecutive iterations without improvement)
- Returned best degree, trained model, and complete CV results dictionary

2. k-Nearest Neighbors Function Development

- Created `train_knn()` function with hyperparameter search
- Implemented cross-validation for k selection (testing k from 1 to max_k)
- Added early stopping mechanism for efficiency
- Returned best k value, trained model, and complete CV results dictionary

3. Hyperparameter Selection

- Used 5-fold cross-validation for both models
- Searched polynomial degrees from 1 to 10
- Searched k values from 1 to 20
- Tracked MSE scores with standard deviations for each hyperparameter

What was done in task [2.1]

4. Model Training

- Trained Polynomial Regression with optimal degree on full dataset
 - Trained k-NN Regression with optimal k on full dataset
 - Generated predictions on training data for both models
-

5. Performance Evaluation

- Calculated training MSE for both models
 - Calculated training cMSE for both models
 - Compared performance against baseline expectations
 - Documented hyperparameter selection results
-

Results and Analysis from task [2.1]

Task [2.2] - Evaluation

What was done in task [2.2]

1. Comprehensive Model Comparison

- Created comparison table with baseline, polynomial regression, and k-NN models
- Included hyperparameter configurations for each model
- Displayed min, max, mean, and standard deviation of errors
- Identified best performing model based on mean cross-validation error

2. Hyperparameter Tuning Visualization

- Plotted polynomial degree vs MSE with confidence intervals
- Plotted k-value vs MSE with confidence intervals
- Marked optimal hyperparameters with vertical lines
- Showed performance trends across hyperparameter ranges

3. Model Predictions Comparison

- Created scatter plots of true vs predicted values for all three models
- Displayed MSE on each plot for direct comparison
- Included ideal prediction line ($y=x$) as reference
- Generated combined and individual visualization plots

What was done in task [2.2]

4. Statistical Analysis

- Computed cross-validation statistics for each model
- Analyzed variance in predictions across folds
- Compared model stability through standard deviation metrics
- Evaluated improvement over baseline model

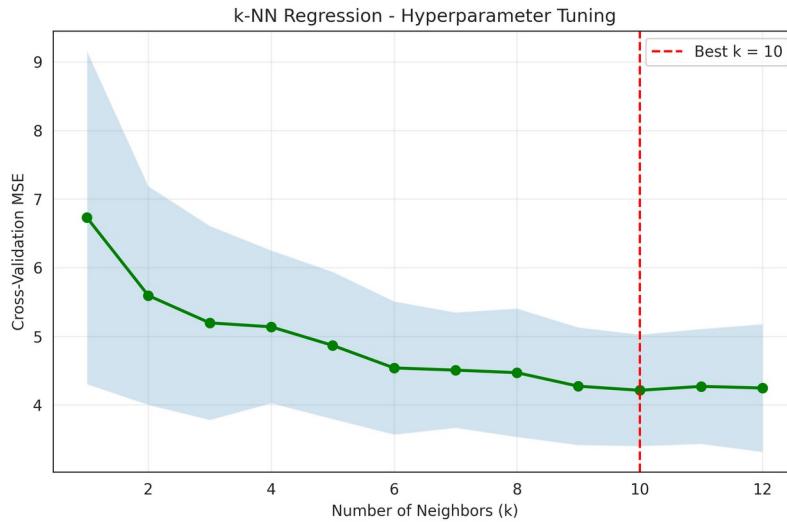
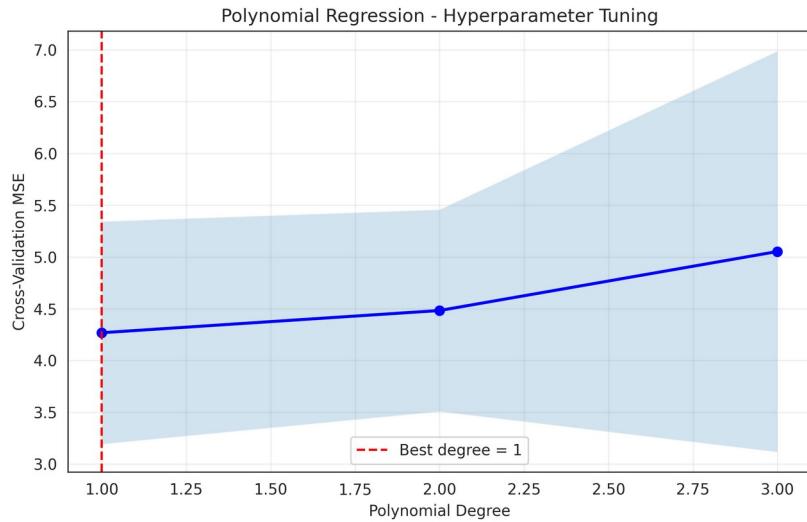
5. Test Set Predictions

- Selected best performing model based on CV results
- Generated predictions for test dataset
- Created submission file for evaluation
- Documented model selection rationale

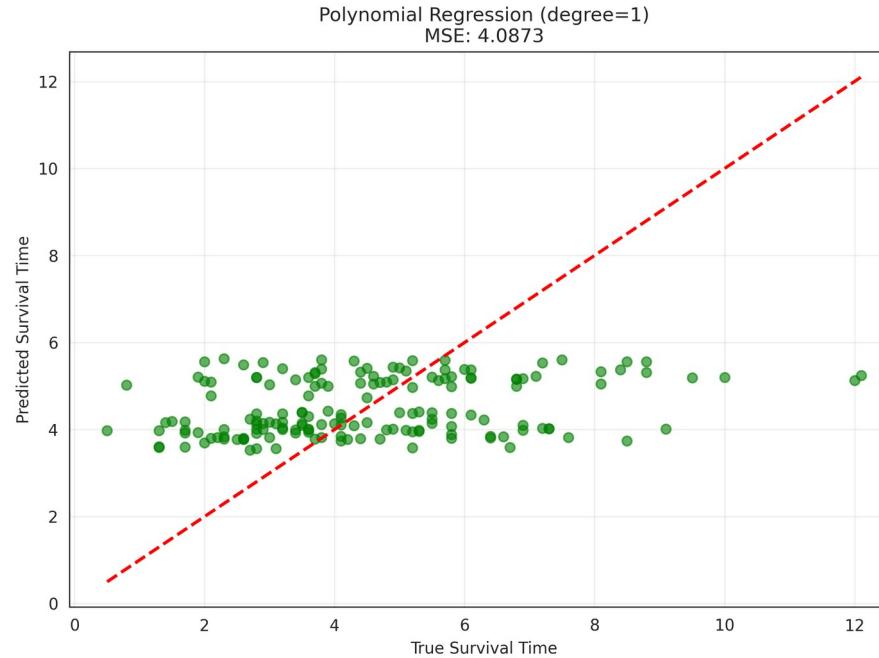
6. Results Documentation

- Saved all comparison plots with task-specific naming
- Generated separate plots for polynomial and k-NN tuning
- Created individual prediction visualizations for each model
- Documented complete evaluation workflow

Results and Analysis from task [2.2]

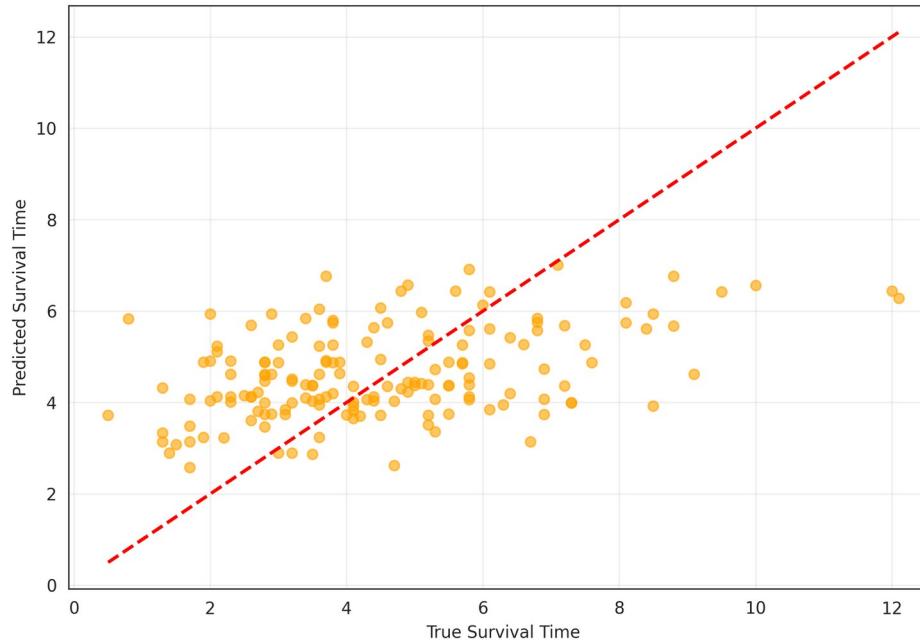


Results and Analysis from task [2.2]

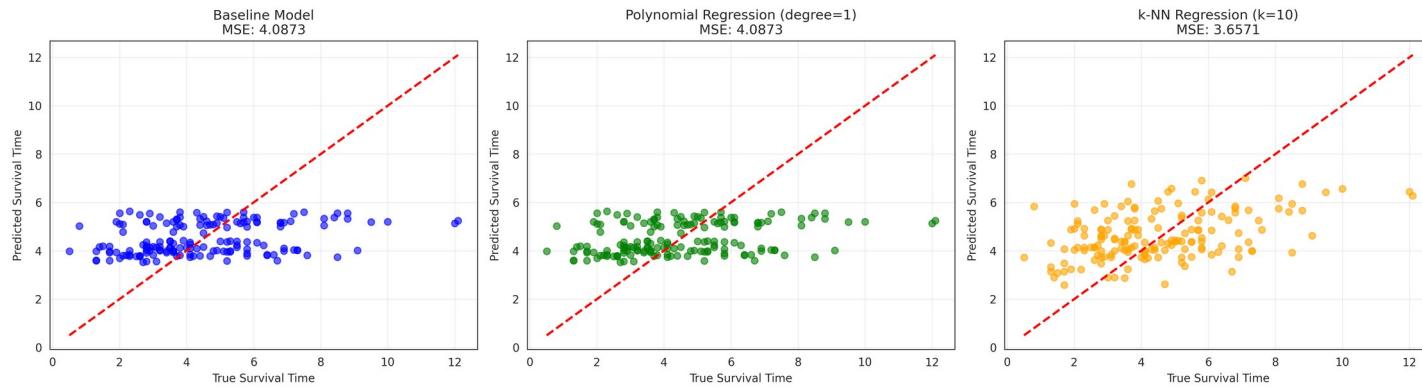


Results and Analysis from task [2.2]

k-NN Regression (k=10)
MSE: 3.6571



Results and Analysis from task [2.2]



Task [3.1] - Missing data imputation

What was done in task [3.1]

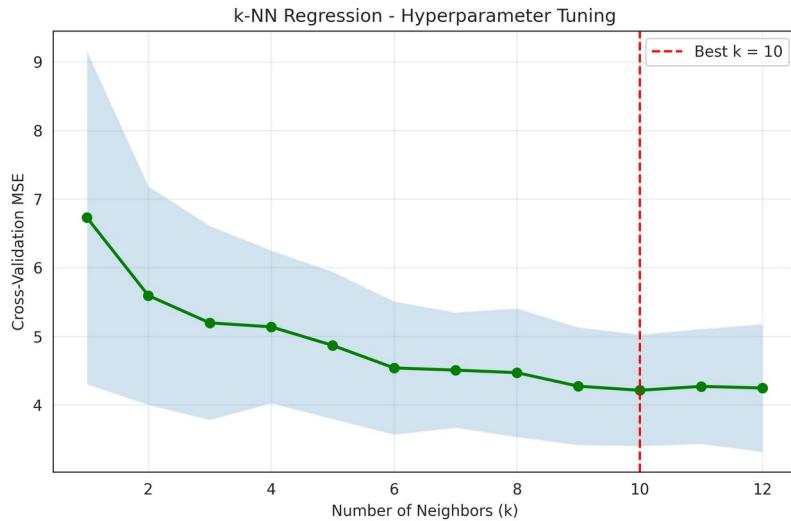
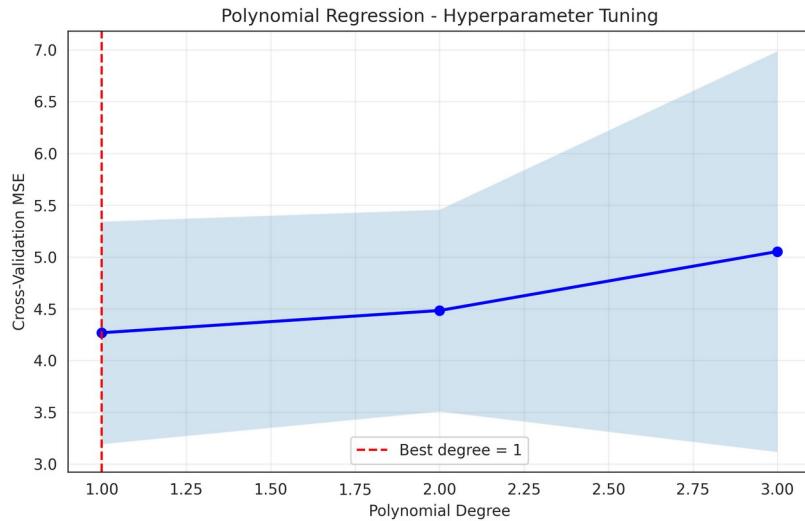
1. Data Preparation

2. Imputation Strategies

3. Model Evaluation

- Load original dataset with missing values
 - Analyze missing value patterns
 - Prepare feature matrix (X) and target variable (y) with missing data intact
-
- Mean Imputation: Replace missing values with column means
 - KNN Imputation: Use k-nearest neighbors to estimate missing values
 - Iterative Imputation: Use Bayesian Ridge regression for multivariate imputation
-
- Train baseline Linear Regression model on each imputed dataset
 - Evaluate using both train/test split and cross-validation approaches
 - Compare performance using cMSE (Censored Mean Squared Error)
 - Test with KNN Regression model for comparison

Results and Analysis from task [3.1]



Task [3.2] - Train models that do not require imputation

What was done in task [3.2]

Results and Analysis from task [3.2]

Task [3.3] - Evaluation

What was done in task [3.3]

1.

Comparison Analysis

- Built comparison table with all strategies: baseline, imputation methods (Mean, KNN, Iterative), and models handling missing data (Decision Tree, HistGradientBoosting, CatBoost AFT)
- Displayed MSE and cMSE metrics for all approaches
- Created y vs y-hat scatter plots for visual comparison of model performance

2.

Combined Approach Testing

- Selected best imputation strategy from Task 3.1: Mean Imputation (cMSE: 1.7645)
- Combined with best model from Task 3.2: CatBoost AFT
- Trained CatBoost AFT on mean-imputed data
- Evaluated performance: cMSE of 3.1236 (worse than native missing handling)

What was done in task [3.3]

- | | |
|----------------------------------|--|
| 3. Best Model Selection | - Compared all strategies including combined approach |
| | - Identified CatBoost AFT with native missing support as best performer (cMSE: 1.7339) |
| 4. Test Predictions & Submission | - Generated predictions on test data using best model (CatBoost AFT) |
| | - Created Kaggle submission file: `handle-missing-submission-xx.csv` |
-

Results and Analysis from task [3.3]

Code Demo

Overall assessment

What went wrong

What went great